

Today's Topics

Data Structures

Common, Useful Commands

Filtering Data

Data Structures

There are different types of data in R, and those types can be organized into different structures in R. The types are

- **Character:** This is text data. For example, the variable “Class” may contain the entries “freshman,” “sophomore,” “junior,” or “senior.”
- **Numeric:** These can be integer or decimal-valued numbers. These are the same as “doubles.”
- **Integer:** These are whole number values, either positive or negative.
- **Logical:** These are either “TRUE” or “FALSE” designations.
- **Complex:** Values with real and imaginary parts, denoted for example as `1+4i`.
- **Date/Time:** These are dates, without or without a specific time stamp associated with them.

These types can be collected and represented as follows:

- **Factor:** A whole column of character data is designated a factor where each unique outcome is called a “level.”
- **Vector:** A sequence of values of the same data type.
- **Matrix:** A two-dimensional (rows and columns) set of values and all must be numeric.
- **Array:** A multi-dimensional set of values all of the same type. A matrix is a special 2-D case.
- **Data Frame:** A very commonly used data structure in R. It is two-dimensional with rows and columns, but each column can be a different data type.
- **List:** A list has slots, and within each slot, you can have a different structure of different sizes stored.

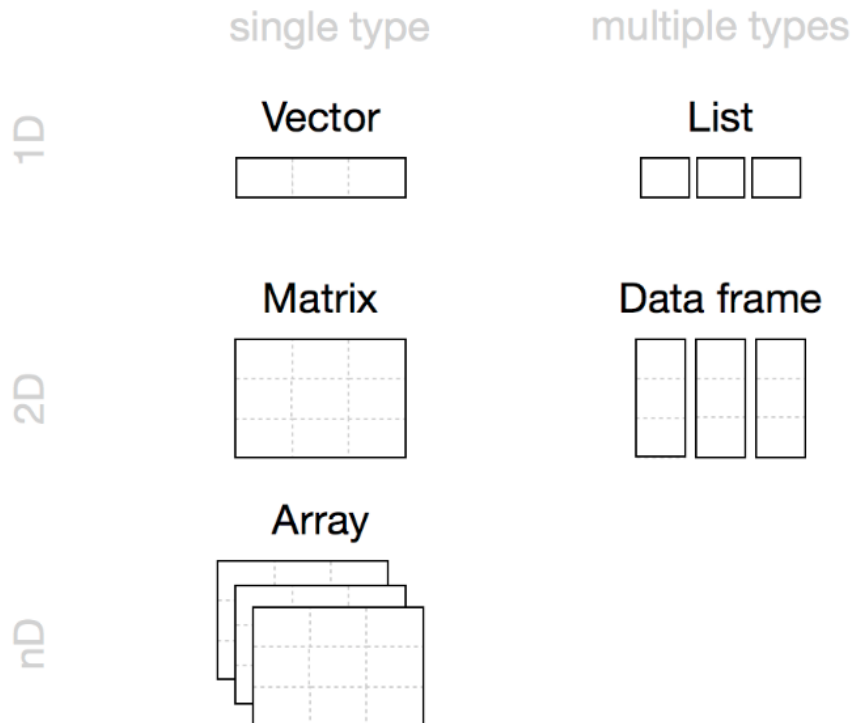


Figure from <https://rstudio-education.github.io/hopr/r-objects.html>

We will work with many data frames, so they deserve some special attention. You can read a dataset into R that is automatically treated as a data frame, or you can create one from scratch, as follows:

```
friends <- data.frame(
  id_number = c(1, 2, 3, 4, 5, 6),
  name = c("Ross", "Rachel", "Monica", "Chandler", "Joey", "Phoebe"),
  birthday = as.Date(c("1966-11-02", "1969-02-11", "1964-06-15", "1969-08-19", "1967-07-25", "1963-07-30")),
  num_daily_coffees = c(3, 5, 2, 3, 2, 8),
  stringsAsFactors = FALSE
)
friends
#   id_number      name  birthday num_daily_coffees
# 1         1      Ross 1966-11-02                 3
# 2         2    Rachel 1969-02-11                 5
# 3         3   Monica 1964-06-15                 2
# 4         4 Chandler 1969-08-19                 3
# 5         5     Joey 1967-07-25                 2
# 6         6   Phoebe 1963-07-30                 8
#friends$name; friends$birthday
```

You extract a particular column from a data frame by using the data frame's name followed by the `$` symbol and then the column name. You can also add a new row or column as

follows:

```
#Add a new column
num_breakups <- c(5, 8, 7, 6, 3, 8)
friends <- cbind(friends, num_breakups)

#Add a new row--requires creating a new data frame for the row
new_person <- data.frame(id_number = 7,
                          name = "Mike",
                          birthday = as.Date("1969-04-06", format="%Y-%m-%d"),
                          num_daily_coffees = 4,
                          num_breakups = 0)
friends <- rbind(friends, new_person)
```

Example: Colorado Covid Data The following data is downloaded from the Colorado Department of Public Health and Environment¹. It contains four main variables:

- The date
- The particular utility
- SARS CoV2 copies of RNA (measured as RNA/liter of water) in wastewater
- The number of new Covid-19 cases

SARS-CoV-2 is the virus that causes COVID-19, and RNA is the genetic material in each copy of the virus. SARS-CoV-2 copies per liter is one measure of how much of the virus is in the wastewater, expressed as a concentration. Studies have shown that individuals who develop COVID-19 often shed detectable SARS-CoV-2 RNA from their systems before, during, and after their infection, so higher levels of SARS-CoV-2 RNA can indicate a rise in cases in a community. Many universities used this method to monitor the wastewater from residence halls to obtain an early warning of a disease outbreak.

```
covid <- read.csv(file="dat/CDPHE_COVID19_Wastewater_Dashboard_Data.csv", header=T)
class(covid)
# [1] "data.frame"
head(covid)
#           Date           Utility SARS_CoV_2_copies_L
# 1 08/15/2020 Metro Wastewater RWHTF - PRC           NA
# 2 08/11/2020           Broomfield           NA
# 3 08/15/2020           Northglenn           NA
# 4 08/11/2020      CO Springs - JD Phillips           NA
# 5 08/11/2020      CO Springs - Las Vegas           NA
# 6 08/15/2020           Pueblo           NA
# Number_of_New_COVID19_Cases_by_ ObjectId
# 1              36              1
# 2              0              2
```

¹<https://cdphe.maps.arcgis.com/apps/opsdashboard/index.html#/d79cf93c3938470ca4bcc4823328946b>

```
# 3      0      3
# 4      6      4
# 5     22      5
# 6      5      6
dim(covid)
# [1] 3498  5
```

The Kindergarten Commands of R

These commonly used commands are so widely used that they are often taken for granted!

- `rm(list=ls())`: removes everything in the global data environment, like erasing a chalkboard and starting over
- `length()`: returns the length of a vector or the number of elements in a matrix. See also, `dim()`, `nrow()`, and `ncol`.
- `c()`: combines values into a group, but they must all be the same data type

```
c(5, 10, -2.5)
# [1] 5.0 10.0 -2.5
```

- `rep()`: repeats the same value (or sequence of values) a certain number of times

```
rep(0, 5)
# [1] 0 0 0 0 0
rep("Hi", len = 3)
# [1] "Hi" "Hi" "Hi"
rep(20:25, times = 2)
# [1] 20 21 22 23 24 25 20 21 22 23 24 25
rep(20:25, each = 2)
# [1] 20 20 21 21 22 22 23 23 24 24 25 25
```

- `seq()`: creates a sequence of numbers from a lower bound to an upper bound, of a given length or separated by a given distance

```
seq(1, 10, len = 21)
# [1] 1.00 1.45 1.90 2.35 2.80 3.25 3.70 4.15 4.60 5.05 5.50 5.95
# [13] 6.40 6.85 7.30 7.75 8.20 8.65 9.10 9.55 10.00
seq(1, 10, by = 0.5)
# [1] 1.0 1.5 2.0 2.5 3.0 3.5 4.0 4.5 5.0 5.5 6.0 6.5 7.0 7.5 8.0
# [16] 8.5 9.0 9.5 10.0
```

- The colon `:` used to rapidly create a sequence of integers

```
1:7
# [1] 1 2 3 4 5 6 7
-3:5
# [1] -3 -2 -1 0 1 2 3 4 5
```

- `unique()` and `table()`:

`unique` will return a list of the unique values in a vector or array and will discard the duplicates.

`table` will tabulate and count up how many times a particular unique element occurs in a vector.

```
#How many utilities are represented in the Covid dataset?
#unique(covid$Utility)
#table(covid$Utility)
```

Filtering Data

Sometimes you want to select only those observations (rows) from a dataset that meet a certain criteria, such as only one particular utility in the Covid data. To do so, you need to know the comparison expressions used in R:

- `x == y` indicates x must equal y
- `x != y` indicates x must NOT equal y
- `x >= y` indicates x must be greater than or equal to y
- `x <= y` indicates x must be less than or equal to y
- `x > y` indicates x must be strictly greater than y
- `x < y` indicates x must be strictly less than y

These are logicals that can be used with the comparisons above:

- `!a` indicates not a
- `a & b` indicates both a AND b must be true
- `a | b` indicates either a OR b must be true

Finally, to apply basic filtering, you need to combine a logical statement with the square brackets. Square brackets allow you to select either the rows or columns in a data frame, indicated as `dataframe[rows, columns]`.

```
#Select row 4 and column 3
friends[4, 5]
# [1] 6

#Selects row 2
friends[2, ]
#   id_number  name  birthday num_daily_coffees num_breakups
# 2          2 Rachel 1969-02-11              5              8
```

```

#Selects column 3
friends[,3]
# [1] "1966-11-02" "1969-02-11" "1964-06-15" "1969-08-19" "1967-07-25"
# [6] "1963-07-30" "1969-04-06"

#Selects id_numbers 1 to 4
friends[1:4, ]
#   id_number      name    birthday num_daily_coffees num_breakups
# 1          1      Ross 1966-11-02              3             5
# 2          2    Rachel 1969-02-11              5             8
# 3          3   Monica 1964-06-15              2             7
# 4          4 Chandler 1969-08-19              3             6

#Selects rows with number of daily coffees over 5
friends[friends$num_daily_coffees > 5, ]
#   id_number      name    birthday num_daily_coffees num_breakups
# 6          6 Phoebe 1963-07-30              8             8

#Selects rows with number of daily coffees over 3 AND number of breakups greater than 3
friends[(friends$num_daily_coffees > 4) & (friends$num_breakups >= 7), ]
#   id_number      name    birthday num_daily_coffees num_breakups
# 2          2    Rachel 1969-02-11              5             8
# 6          6 Phoebe 1963-07-30              8             8

#Selects rows with number of daily coffees over 3 OR number of breakups greater than 3
friends[(friends$num_daily_coffees > 4) | (friends$num_breakups >= 7), ]
#   id_number      name    birthday num_daily_coffees num_breakups
# 2          2    Rachel 1969-02-11              5             8
# 3          3   Monica 1964-06-15              2             7
# 6          6 Phoebe 1963-07-30              8             8

```

Example: Colorado Covid Data Find the following subsets:

1. Find those rows of the dataset that pertain only to the Boulder utility.
2. Select those observations whose new Covid-19 cases are greater than 200. What proportion of the dataset is this?
3. Select the rows that do not have any NAs in the RNA column.

```

#Question 1
boulder_covid <- covid[covid$Utility == "Boulder", ]
#boulder_covid

#Question 2

```

```

high_counts <- covid[covid$Number_of_New_COVID19_Cases_by_ > 200, ]
nrow(high_counts) / nrow(covid) *100
# [1] 5.517439

#Question 3
covid_complete <- covid[is.na(covid$SARS_CoV_2_copies_L) == FALSE, ]
covid_complete <- covid[!is.na(covid$SARS_CoV_2_copies_L), ]
#Can also be split into two steps.
#rna_present_index <- is.na(covid$SARS_CoV_2_copies_L)
#covid_complete <- covid[!rna_present_index, ]

```

Are any of the columns of this dataset redundant? How could you quickly remove a column?

```

covid <- covid[ , -5]

```